

Stock Discovery Plugin — Walkthrough

A research copilot for fundamental discovery on Indian markets (NSE/BSE)

What this is

A Claude Code plugin built around six skills that mirror how a fundamentals-driven investor actually researches a stock. The user picks a sector and growth thesis; the plugin translates that into companies, pulls 10-12 years of financials per ticker, surfaces anomalies in the balance sheet / P&L / cash flow, lets the user accept/reject each candidate with a written reason, and learns from those reasons so every future research run gets sharper.

The user remains the decision-maker. The plugin is explicitly not a stock picker, not a screener, and not a technical-analysis tool — it brings inputs, the user applies subjectivity.

The six skills

Skill	What it does
thesis-to-products	Translate a sector + growth thesis into specific products / sub-themes.
find-companies	Surface NSE/BSE tickers — primary, auxiliary, adjacent — for a sub-theme.
anomaly-scan	Pull 10-12yr financials from Screener.in, surface anomalies vs the user's criteria.
compare	Free-form cross-stock Q&A across the shortlist — callable any time, mid-flow.
decide	Capture accept/reject + reason; cross-question when reasoning is thin.
philosophy-refresh	Distil all past decisions into a prescriptive philosophy file used by every other skill.

How context flows

Every skill is short-lived (one chat turn). Persistence happens through markdown files all skills read on entry and write on exit:

state/shortlist.md — the live working set of tickers and their statuses (*candidate* / *scanned* / *accepted* / *rejected*). Hand-editable.

memory/criteria.md — anomaly checklist used by anomaly-scan and compare.

memory/philosophy.md — distilled judgment, loaded into every qualitative skill.

decisions/<TICKER>/<date>.md — full audit trail of accept/reject reasoning.

Because there is no in-memory state — only files — any skill can be called at any point in the workflow and pick up exactly where things stand. This is what makes *compare* safe to call mid-flow without breaking continuity.

The walkthrough

What follows is the actual experience your friend would have when he opens Claude Code in the plugin directory. The blue indented blocks are **what your friend types**. The rest is what the plugin produces in response. All financial numbers shown are real, pulled live from Screener.in during this demo and committed to the cache.

Theme used in this walkthrough: **Indian railway modernisation post the FY26 budget**.

Turn 1 — Cold start (thesis-to-products)

***Friend types:** I'm bullish on Indian railway modernisation post the FY26 budget. Vande Bharat ramp, Kavach rollout, station redevelopment. Help me find stocks.*

That phrasing trips the **thesis-to-products** skill. It reads **memory/philosophy.md** — empty, fresh repo — and replies asking for two missing inputs before it commits to a list:

Plugin replies: Running fresh, no philosophy yet. Two clarifications before I break this down:

1. Time horizon — 1-2yr swing or 3-5yr hold? (Affects whether I lean speculative on smaller-cap pure plays.)
2. Are PSUs (RVNL, IRCON, IRFC) in scope, or only private compounders?

Quick read of the thesis itself: railway capex outlay was up ~5% in the FY26 budget, Kavach had **■1,100cr** earmarked, Vande Bharat sleeper variants in tender. The thesis is real but the multiples on rolling-stock names already price in a lot. Worth tagging in your head.

***Friend types:** 2-3 year hold. Both PSU and private fine.*

Plugin replies:

Sub-theme	Why it captures the thesis	Timing	Confidence
Rolling stock (Vande Bharat coaches, metro, freight wagons)	Direct beneficiary of coach orders + Vande Bharat sleeper tenders	1-3 yrs	high
Signalling & safety (Kavach, train-control electronics)	■1,100cr FY26 outlay; equipment supplier shortlist tight	2-3 yrs	high
Railway electrification & T&D	100% electrification target; multi-year EPC pipeline	2-4 yrs	high
Wagons (private freight + DFC)	DFC ramp + private wagon-leasing scheme	1-3 yrs	medium
Railway PSUs (project execution, financing)	RVNL/IRCON/IRFC — operating leverage on order book	2-4 yrs	medium
Station redevelopment (capex, real-estate adjacent)	1,200+ stations under redevelopment, lumpy revenue	3-5 yrs	low

What I'd push back on: Rolling-stock multiples are stretched — Titagarh and JWL trade at 50-70x. If your edge is just "the orders are real", consensus already knows. The under-priced angle may be in Kavach electronics or T&D where the order book hasn't fully converted yet. Worth thinking about before you pick a sub-theme.

Turn 2 — Picking a sub-theme (find-companies)

Friend types: Let's go with rolling stock. Find me companies.

find-companies fires. Reads **state/shortlist.md** (empty), **memory/philosophy.md** (empty), runs grounded research, returns three tables, and writes the rows to disk.

Primary players

Ticker	Company	Why included	Conf.
TITAGARH.NS	Titagarh Rail Systems	Vande Bharat consortium, freight wagons, propulsion JV	high
JWL.NS	Jupiter Wagons	Wagons + Vande Bharat propulsion JV with Stadler	high
TEXRAIL.NS	Texmaco Rail	Wagons + EPC, tie-up with Wabtec	high
BEML.NS	BEML Ltd	Metro coaches + Vande Bharat consortium	high

Auxiliary

Ticker	Company	Auxiliary role	Conf.
RAMKRISHNA.NS	Ramkrishna Forgings	Wheels + axles for Indian Railways	high
STELHOLD.NS	Steel Strips Wheels	Wheels for rolling stock	medium

Adjacent

Ticker	Company	Adjacency	Conf.
RVNL.NS	Rail Vikas Nigam	EPC + wagon leasing; not pure rolling-stock but adjacent	high
IRFC.NS	IR Finance Corp	Financing arm — operating leverage on rolling-stock orders	medium

Added all 8 to **state/shortlist.md** with status *candidate*. Run *anomaly-scan --all-candidates* next, or call out specific ones to scan first.

Turn 3 — Anomaly scan on a real ticker

Friend types: Scan TITAGARH first. I want to see how it actually reads.

anomaly-scan fires. Reads **memory/criteria.md**, then runs `python tools/fetch_screener.py TITAGARH.NS`, then applies the criteria to the data and produces a report. The numbers below were pulled live from Screener during this demo.

Headline summary

A genuine turnaround story, not a smooth compounder. Loss-making for 9 of 10 years between FY14 and FY22 (cumulative net loss ~■240cr). FY23 inflection: ■126cr profit. FY24-FY25 sustained at ■275-286cr. Revenue scaled 6× (■624cr → ■3,868cr) but the inflection is recent. The story is real, the question is whether you trust the durability.

P&L anomalies

Pattern	Years	Magnitude	Severity
Decade of losses then sudden turnaround	FY14-FY22 → FY23	-■147cr in FY18 to +■126cr in FY23	flag (this is the thesis)
OPM expansion 1% → 12%	FY14 → FY24	1100bps, mostly in FY22-FY24	flag — durability test
"Other Income" volatile, sometimes negative	FY20 (-■60cr), FY22 (-■61cr)	Suggests forex / exceptional impacts	concern
Tax rate erratic (-59% to 237%)	FY18-FY21	Loss-period normalisation	noise

Balance Sheet anomalies

Pattern	Years	Magnitude	Severity
Borrowings spike then crash	FY18 → FY22 → FY23	■624cr → ■920cr → ■353cr; de-leveraged via QIP	flag
Reserves jump from ■940cr → ■2,191cr	FY23 → FY24	~■1,250cr equity raise (QIP confirmation)	concern — dilution check
Equity capital bumped 24 → 27cr	FY24	Confirms the QIP	noise
CWIP spike to ■174cr	FY24	New capacity in flight; ■128cr by FY25	flag — watch revenue lift
Borrowings re-rising FY24-FY25	FY24 → FY25	■166cr → ■637cr; re-leveraging despite recent equity raise	concern

Cross-statement patterns

The turnaround is real but capital-intensive. They raised ■1,250cr in equity in FY24 and re-leveraged borrowings 4× in two years. The thesis works only if the new capacity converts to revenue at margin.

Equity dilution before re-rating is unusual — usually management waits for higher prices. Either they expect bigger capex needs ahead, or they doubt durability of multiples.

9 years of losses means the entire institutional memory of this stock pre-FY23 is bad. New investors are pricing the *new* Titagarh; you're betting that old-era issues don't recur.

What to investigate further

1. What's the FY24 capex actually building — VB sleeper coaches? Propulsion? Wagon capacity? Match it to the order book.
2. Pull the CF statement detail and check whether the 5-yr OCF/PAT ratio supports the turnaround narrative or smells like accruals.
3. Promoter holding trend post-QIP — dilution often shifts the cap table.
4. Comparable: JWL.NS scaled from a tiny base too. Cleaner story, similar ramp?

Turn 4 — Comparative question, no upstream re-run (compare)

Friend types: Wait — before I keep going, give me debt-to-equity, 4-yr revenue CAGR, FY25 OPM, and equity dilution since FY21 across the 3 I'd actually consider: TITAGARH, JWL, RVNL.

compare fires. Reads `state/shortlist.md`, sees all three are in `.cache/` already (no re-fetch needed), pulls the lines, computes the metrics:

Ticker	Company	Sales FY25 (cr)	4yr Rev CAGR	OPM FY25	D/E FY25	Equity dilution FY21→25
JWL.NS	Jupiter Wagons	3,963	41.2%	14%	0.18	+9.6%
TITAGARH.NS	Titagarh Rail	3,868	26.2%	11%	0.25	+12.5%
RVNL.NS	Rail Vikas Nigam	19,923	6.6%	6%	0.57	0%

Observation

JWL is the cleanest growth story — highest revenue CAGR, highest margin, lowest leverage, smallest dilution. Trade-off: only 5 years of public history (listed ~2020), so you have less to anchor on.

Titagarh has the bigger absolute turnaround but came from a worse place — 9 years of losses pre-FY23 and the FY24 QIP dilution is real.

RVNL is a different animal — 5× the size, 6% margins, no dilution but ~3× the leverage. It's an EPC PSU not a manufacturer; comparing it to JWL/Titagarh on margin is apples-to-oranges. Worth keeping for diversification but the thesis here is operating leverage on order-book conversion, not margin expansion.

Notice: no status changes, no fetches, no writes. The friend can call this 10 more times with different metric combinations and it costs nothing — that's the cache + shared-state design paying off.

Turn 5 — Thin reject reason → cross-question (decide)

Friend types: Reject Titagarh. Looks weak.

decide fires. Reads `state/shortlist.md` (sees TITAGARH at status *scanned*), peeks at `.cache/TITAGARH.NS.json`, sees the reason is two words — and gates the save.

Plugin replies: Holding off on saving — "looks weak" doesn't tell future-you (or me) anything actionable. Two follow-ups:

1. The data shows a real turnaround (■126cr → ■275cr profit in 2 yrs) and 26% revenue CAGR. What does "weak" mean to you here — the ■1,250cr QIP dilution, the re-leveraging despite the equity raise, doubts about durability, or something else?

2. You're keeping JWL on the table even though it's a similar story. What disqualifies Titagarh that JWL passes? If the answer is "JWL is just better on the metrics" that's fine — say it that way and I'll save the comparative reason.

Friend types: *Both. The QIP + immediate re-leverage tells me management isn't confident in the cash flow durability. They diluted and added debt — that's belt-and-suspenders. JWL is similar size but D/E 0.18 vs 0.25 and dilution was smaller. So it's a relative call — Titagarh's capital structure choices made me trust them less than JWL on the same thesis.*

Now **decide** saves the verdict — with the substantive reason, the anomalies that drove it, and the cross-questioning Q&A — to **decisions/TITAGARH.NS/2026-05-10.md**. Updates **state/shortlist.md** → status *rejected*. Triggers **philosophy-refresh**.

Turn 6 — Philosophy-refresh distils what just happened

philosophy-refresh reads every file in **decisions/**, finds 1 decision (TITAGARH reject), distils it into a fresh philosophy file. The output is prescriptive, not descriptive — phrased so it can be applied as a lens on future runs:

Strong negatives

Belt-and-suspenders funding — when a company raises fresh equity and *also* re-leverages debt within 1-2 years on the same capex story, treat it as a signal that management itself isn't confident in operating cash flow durability. Titagarh raised ■1,250cr (FY24 QIP) and re-leveraged borrowings 4× (■166cr → ■637cr) over FY24-FY25. Reject — even though the operating turnaround was real.

Recent dilution before the operating story is durable. If the turnaround is <2 years old and management has already diluted equity, the read is that they expect the next chapter to need more capital — fine in itself, but it shifts the per-share math underneath you.

Subjectivity dimensions

Capital structure as a tell on management conviction. Titagarh framed: "they diluted *and* added debt — that's belt-and-suspenders." Use this lens any time a company runs both equity raise and debt expansion in adjacent years on a stated growth capex story.

Relative ranking on the same thesis. Where two companies express the same investment thesis (Titagarh and JWL on rolling stock), the user prefers to reject the weaker capital structure even if both have viable businesses. The decision is comparative, not absolute.

Turn 7 — The compounding effect

A week later, your friend opens Claude Code in the same folder with a new theme.

***Friend types:** Find me companies in Indian capital goods playing the power-T&D capex cycle.*

find-companies fires. Reads **state/shortlist.md** (8 railway tickers — irrelevant), reads **memory/philosophy.md** (now non-empty). Notices the user's belt-and-suspenders / capital-structure lens. Returns the table — and *automatically annotates* which suggestions might trip his lens:

Plugin replies: Loaded your philosophy: capital structure is a primary lens, you penalise belt-and-suspenders funding, and you rank comparatively on the same thesis.

Ticker	Company	Why included	Capital-structure flag (per your lens)
KEC.NS	KEC International	Power T&D EPC, global presence	Stable cap structure, no recent QIP — passes
KALPATPOWR.NS	Kalpataru Projects	T&D + buildings + B&F	Recent QIP in FY24 — flag for capital-structure review
TRANSWORLD.NS	Transformers & Rectifiers	T&D equipment, FY24 inflection	QIP + warrants in FY24 — flag, similar pattern to Titagarh

That's the compounding loop. One decision changed how every future search runs, without the user re-explaining himself. The plugin doesn't disqualify the flagged tickers — that's still his call — it just surfaces the pattern early so the friend can decide where to spend his investigation time.

What's actually on disk after this session

Everything below was committed by the plugin during this walkthrough — these are real files in the repo, not screenshots.

File	What it contains
state/shortlist.md	8 railway tickers; TITAGARH flipped to <i>rejected</i> , others remain <i>candidate</i> .
.cache/TITAGARH.NS.json	12 yrs of P&L / BS / CF / Ratios; ~12kB; 7-day TTL.
.cache/JWL.NS.json	5 yrs of statements; cached for compare to reuse.
.cache/RVNL.NS.json	11 yrs of statements; cached for compare to reuse.
decisions/TITAGARH.NS/2026-05-10.md	Full reject record: reason, anomalies, cross-questioning Q&A, open follow-ups.
memory/philosophy.md	Distilled philosophy after 1 decision — capital-structure lens, comparative-ranking lens.

Why this design

The user never types the same lens twice. The QIP-plus-re-leverage frame invented in turn 5 is now permanent context. Every future *find-companies*, *anomaly-scan*, and *compare* run will surface it as a lens before the user has to remember to ask.

Mid-flow questions don't break flow. The walkthrough went sector → companies → scan → comparative metric → verdict in one chat, not five tools. The *compare* step took ~1 second because Titagarh, JWL, and RVNL were already cached from the earlier *anomaly-scan*.

The cross-question wasn't scripted. It referenced the actual ■126cr → ■275cr profit data from the cache, and called out that the friend was keeping JWL on the same thesis. That comparative pressure is what forced him to articulate "it's a relative call" — turning a vibe into a reusable rule.

The philosophy file is short and prescriptive on purpose. Long philosophy = bloated context on every future skill run. Brevity is a feature, not a shortcut.

Generated 2026-05-10 from the live demo session. Tickers, financial data, and anomaly observations are real and pulled from Screener.in. The decision and philosophy text are synthesised for illustration — your friend's actual usage would produce his own.